Apache Solr crash course

Tommaso Teofili

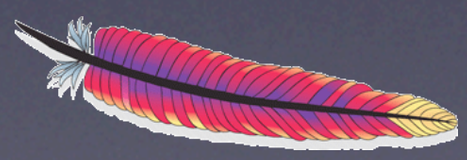

Agenda

- IR
- Solr
- Tips&tricks
- Case study
- Extras

Information Retrieval

- “Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers)” - P. Nayak, Stanford University

Inverted index

- Each document has an id and a list of terms
- For each term t we must store a list of all documents that contain t
- Identify each document by its id

	A	B
1	term	docs
2	pizza	3, 5
3	solr	2
4	lucene	2, 3
5	sourcesense	2, 4
6	paris	1, 10
7	tomorrow	1, 2, 4, 10
8	caffè	3, 5
9	big	6
10	brown	6
11	fox	6
12	jump	6
13	the	1, 2, 4, 5, 6, 8, 9

IR Metrics

- How good is an IR system?
- Precision: Fraction of retrieved docs that are relevant to user's information need
- Recall: Fraction of relevant docs in collection that are retrieved

Apache Lucene

- Information Retrieval library
- Inverted index of documents
- Vector space model
- Advanced search options (synonims, stopwords, similarity, proximity)

Lucene API - indexing

- Lucene indexes are built on a Directory
- Directory can be accessed by IndexReaders and IndexWriters
- IndexSearchers are built on top of Directories and IndexReaders
- IndexWriters can write Documents inside the index
- Documents are made of Fields
- Fields have values
- Directory > IndexReader/Writer > Document > Field

Lucene API - searching

- Open an IndexSearcher on top of an IndexReader over a Directory
- Many query types: TermQuery, MultiTermQuery, BooleanQuery, WildcardQuery, PhraseQuery, PrefixQuery, MultiPhraseQuery, FuzzyQuery, NumericRangeQuery, ...
- Get results from a TopDocs object

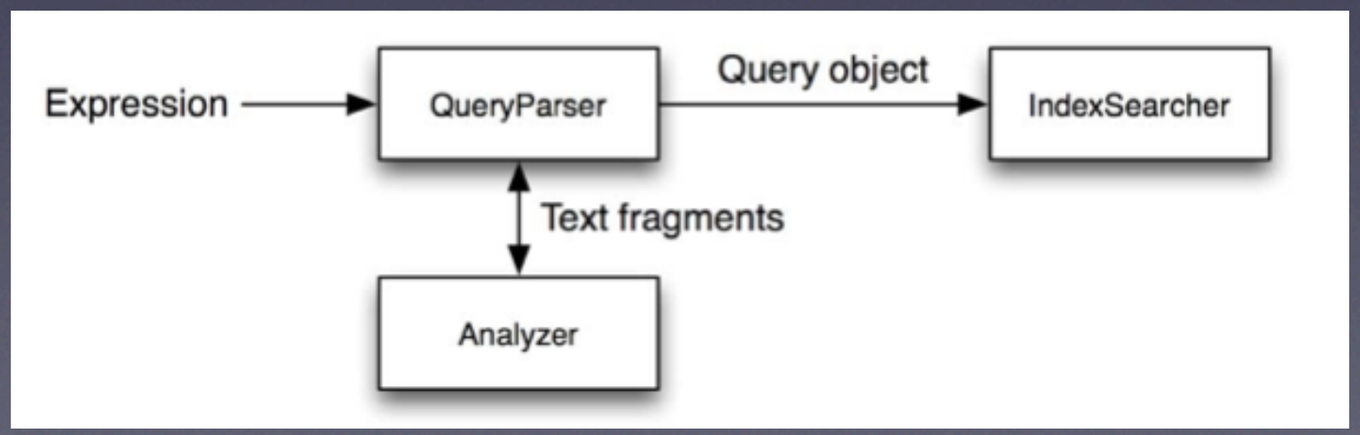

Apache Solr

- Ready to use enterprise search server
- REST (and programmatic) API
- Results in XML, JSON, PHP, Ruby, etc...
- Exploit Lucene power
- Scaling capabilities (replication, distributed search, ...)
- Administration interface
- Customizable via plugins

Solr Admin

Apache Solr

Examples: Simple Spatial Group By

Find: Invia Ripristina

Boost by Price

Field Facets

cat

- Electronics (14)
- Memory (3)
- Connector (2)
- Graphics Card (2)
- Hard Drive (2)
- Monitor (2)
- Search (2)
- Software (2)
- Camera (1)
- Copier (1)
- Multifunction Printer (1)
- Music (1)
- Printer (1)
- Scanner (1)

manu_exact

- Apache Software Foundation (2)
- Belkin (2)
- Canon Inc. (2)

28 results found in 42 ms Page 1 of 3

Test with some GB18030 encoded characters More Like This

Price: € 0,00

Features: No accents here 这是一个功能 This is a feature (translated) 这份文件是很有光泽 This document is very shiny (translated)

In Stock: true

Samsung SpinPoint P120 SP2514N - hard drive - 250 GB - ATA-133 More Like This

Price: € 92,00

Features: 7200RPM, 8MB cache, IDE Ultra ATA-133 NoiseGuard, SilentSeek technology, Fluid Dynamic Bearing (FDB) motor

In Stock: true

Larger Map

Maxtor DiamondMax 11 - hard drive - 500 GB - SATA-300 More Like This

Price: € 350,00

Features: SATA 3.0Gb/s, NCQ 8.5ms seek 16MB cache

In Stock: true

Larger Map

Apache Solr 3.1.0

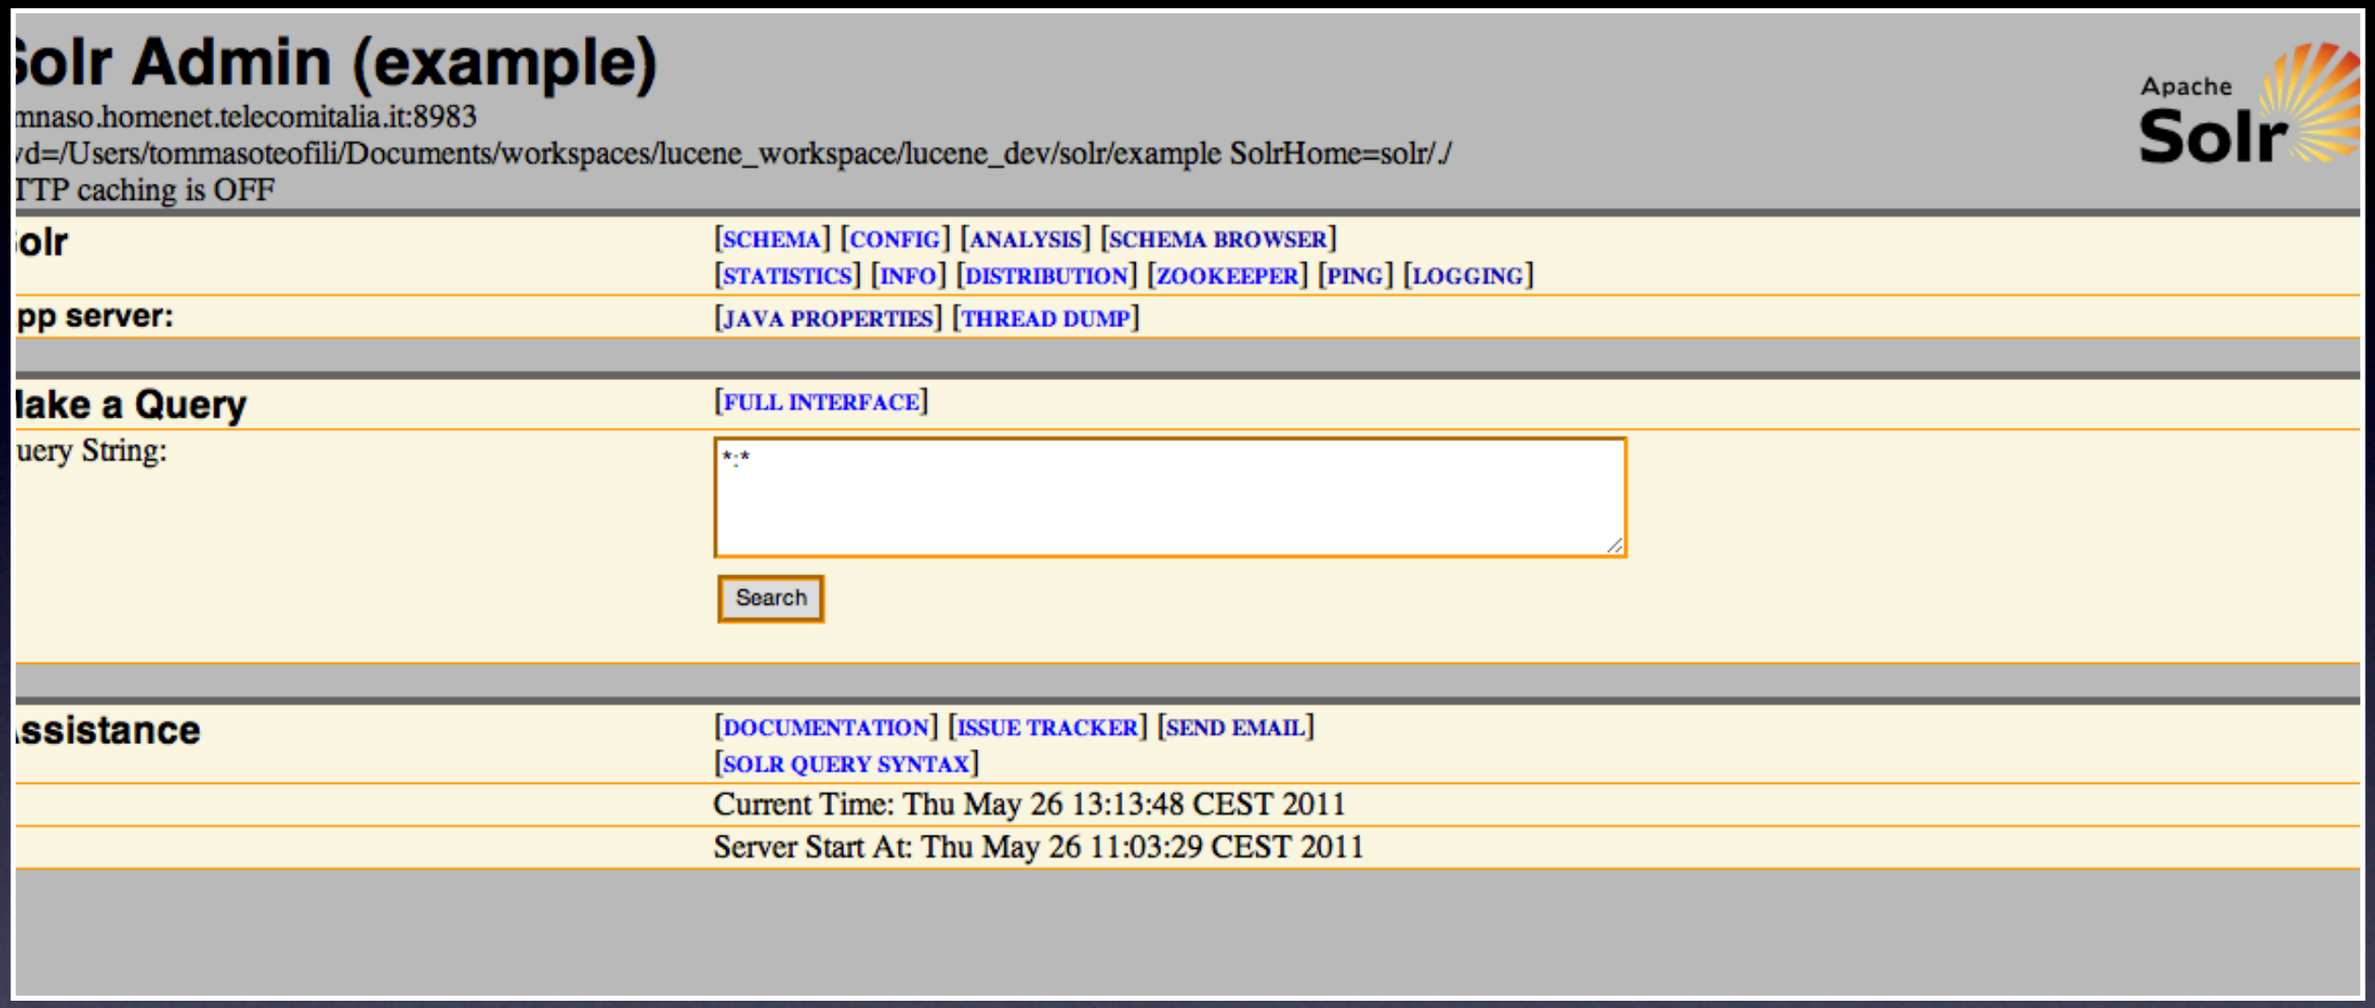

Apache Solr - admin UI

Solr - project status

- Solr 3.1.0 version released in March 2011
- Lucene/Solr is now a single project
- Huge community
- Backed by Lucid Imagination

Solr basic configuration

- schema.xml
  - contains types definitions for field analysis (field type+tokenizers+filters)
  - contains field definitions
- solrconfig.xml
  - contains the Solr instance configuration

Solr - schema.xml

- Types (with index/query Analyzers)
- Fields with name, type and options
- Unique key
- Dynamic fields
- Copy fields

Solr - content analysis

- define documents' model
- each document consists of fields
  - each field
    - has attributes telling Solr how to handle its contents
    - contains free text, keywords, dates, numbers, etc.

Solr - content analysis

- Analyzer: create tokens using a Tokenizer and, eventually, some filters (TokenFilters)
- Each field can define an Analyzer at ‘query’ time and another at ‘index’ time, or the same in both cases
- Each field can be indexed (searchable), stored (possibly fetched with results), multivalued, required, etc.

Solr - content analysis

- Commonly used tokenizers:
  - WhitespaceTokenizerFactory
  - StandardTokenizerFactory
  - KeywordTokenizerFactory
  - PatternTokenizerFactory
  - HTMLStripWhitespaceTokenizerFactory
  - HTMLStripStandardTokenizerFactory

Solr - content analysis

- Commonly used TokenFilters:
  - SnowballPorterFilterFactory
  - StopFilterFactory
  - LengthFilterFactory
  - LowerCaseFilterFactory
  - WordDelimiterFilterFactory
  - SynonymFilterFactory
  - PatternReplaceFilterFactory
  - ReverseWildcardFilterFactory
  - CharFilterFactories (Mapping,HtmlString)

Solr Admin (example)

tomnaso.homenet.telecomitalia.it:8983
cwd=/Users/tommasoteofili/Documents/workspaces/lucene_workspace/lucene_dev/solr/example SolrHome=solr/./
HTTP caching is OFF

Apache Solr

Field Analysis

Field type: text

Field value (Index)
verbose output ☐
highlight matches ☑

Sourcesense, making sense of open source

Field value (Query)
verbose output ☐

Open Source

Analyze

Index Analyzer

Sourcesense,	making	sense	of	open	source
Sourcesense,	making	sense		open	source
Sourcesense	making	sense		open	source
sourcesense	making	sense		open	source
sourcesense	making	sense		open	source
sourcesens	make	sens		open	sourc

Query Analyzer

Open	Source
Open	Source
Open	Source
Open	Source
open	source
open	source
open	sourc

Debugging analysis

Solr - solrconfig.xml

- Data directory (where Solr will write the Lucene index)
- Caches configuration: documents, query results, filters
- Request handlers definition (search/update handlers)
- Update request processor chains definition
- Event listeners (newSearcher, firstSearcher)
- Fine tuning parameters
- ...

Solr - indexing

- Update requests on index are given with XML commands via HTTP POST
- <*add*> to insert and update
  - *<add> <doc boost="2.5">*
    - *<field name="employeeId">05991</field>*
    - *</doc></add>*
- <*delete*> to remove by unique key or query
  - *<delete><id>05991</id></delete>*
  - *<delete><query>office:Bridgewater</query></delete>*
- <*commit/*> reopen readers on the new index version
- <*optimize/*> optimize index internal structure for faster access

Solr - basic indexing

- REST call - XML/JSON
  - curl ‘http://localhost:8983/solr/update?commit=true’ -H "Content-Type: text/xml" --data-binary '<add><doc><field name="id">testdoc</field></doc></add>'
  - curl 'http://localhost:8983/solr/update/json?commit=true' -H 'Content-type:application/json' -d ' { "add": {"doc": {"id" : "TestDoc1", "title" : "test1"} } }’

Solr - binary files indexing

- Many documents are produced in (properietary) binary formats : PDF, RTF, XLS, etc.
- Apache Tika integrated in Solr REST service for indexing such documents
- curl "http://localhost:8983/solr/update/extract?literal.id=doc1&commit=true" -F "myfile=@tutorial.html"

Solr - index analysis

- Luke is a tool for navigating Lucene indexes
- For each field : top terms, distinct terms, terms histogram, etc.
- LukeRequestHandler :
  - http://localhost:8983/solr/admin/luke?wt=xslt&tr=luke.xsl

Solr - data import handler

- DBMS
- FileSystem
- HTTP

Apache Solr

examples: Simple Spatial Group By

Find: SATA hard drive Invia Ripristina

Boost by Price

Field Facets

cat

- Electronics (3)
- Hard Drive (2)
- Copier (1)
- Multifunction Printer (1)
- Printer (1)
- Scanner (1)

manu_exact

- Canon Inc. (1)
- Maxtor Corp. (1)
- Samsung Electronics Co. Ltd. (1)

Query Facets

- GB (2)

Range Facets

3 results found in 21 ms Page 1 of 1

Maxtor DiamondMax 11 - *hard drive* - 500 GB - *SATA*-300 More Like This

Price: € 350,00

Features: *SATA* 3.0Gb/s, NCQ

In Stock: true

Larger Map

Samsung SpinPoint P120 SP2514N - *hard drive* - 250 GB - ATA-133 More Like This

Price: € 92,00

Features: 7200RPM, 8MB cache, IDE Ultra ATA-133 NoiseGuard, SilentSeek technology, Fluid Dynamic Bearing (FDB) motor

In Stock: true

Larger Map

Canon PIXMA MP500 All-In-One Photo Printer More Like This

Price: € 179,99

Features: memory card: CompactFlash, Micro *Drive*, SmartMedia, Memory Stick, Memory Stick Pro, SD Card

In Stock: true

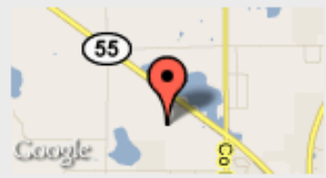

Larger Map

Solr - searching

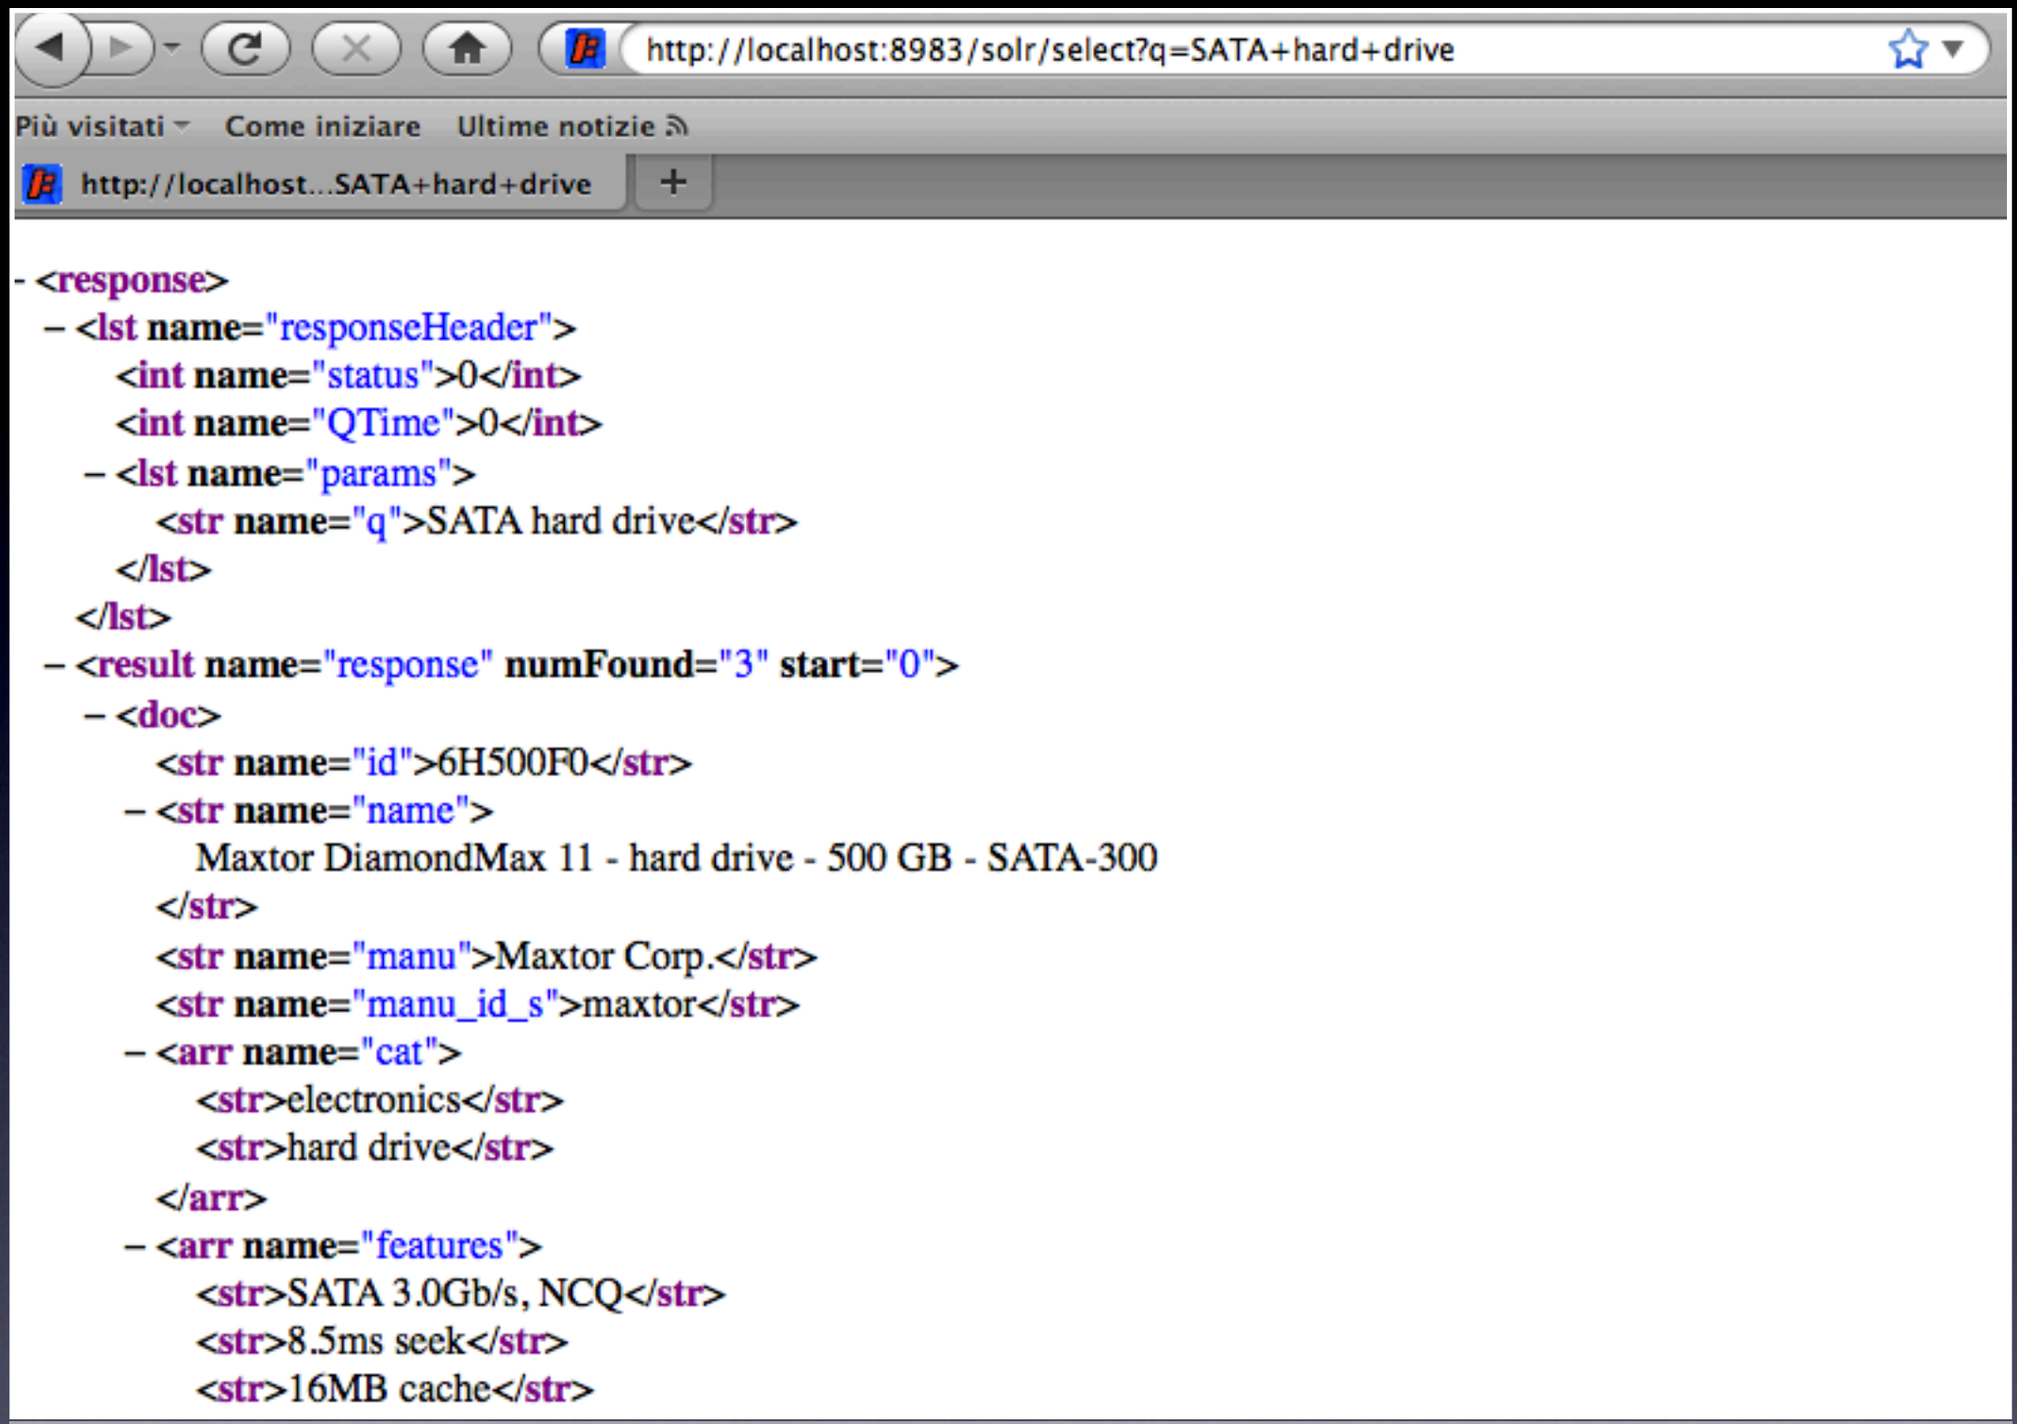

Solr - searching

Solr - query syntax

- query fields with fieldname:value
- + - AND OR NOT operators
- Range queries on date or numeric fields, ex: timestamp:[* TO NOW]
- Boost terms, ex: people^4 profits
- Fuzzy search, ex: roam~0.6
- Proximity search, ex: “apache solr”~2
- ...

Solr - basic search

- parameters:
  - q: the query
  - start: offset of the first result
  - rows: max no. of results returned
  - fl: comma separated list of fields to return
  - defType: specify the query parser
  - debugQuery: enable query debugging
  - wt: result format (xml, json, php, ruby, javabin, etc)

Solr - query parsers

- Most used:
  - Default Lucene query parser
  - DisMax query parser
  - eDisMax query parser

Solr - highlighting

- can be done on fields with stored=”true”
- returns a snippet containing the higlighted terms for each doc
- enabled with hl=true&hl.fl=fieldname1,fieldname2

Solr - sorting results

- Sorting can be done on the "score" of the document, or on any multiValued="false" indexed="true" field provided that field is either non-tokenized (ie: has no Analyzer) or uses an Analyzer that only produces a single term
- add parameter &sort=score desc, inStock desc, price asc
- can sort on function queries (see later)

Solr - filter queries

- get a subset of the index
- place it in a cache
- run queries for such a “filter” in memory
- add parameter &fq=category:hardware
- if multiple fq parameters the query will be run against the intersection of the specified filters

Solr - facets

- facet by:
  - field value
  - arbitrary queries
  - range
- can facet on fields with indexed=”true”

Solr - function queries

- allow deep customization of ranking :
  - http://localhost:8983/solr/select/?fl=score,id&q=DDR&sort=termfreq(text,memory)%20desc
- functions : sum, sub, product , div ,pow, abs, log, sqrt, map, scale, termfreq, ...

Solr - query elevation

- useful for “marketing”
- configure the top results for a given query regardless of the normal Lucene scoring
- http://localhost:8983/solr/elevate?q=best%20product&enableElevation=true

Solr - spellchecking

- collects suggestions about input query
- eventually correct user query with "suggested" terms

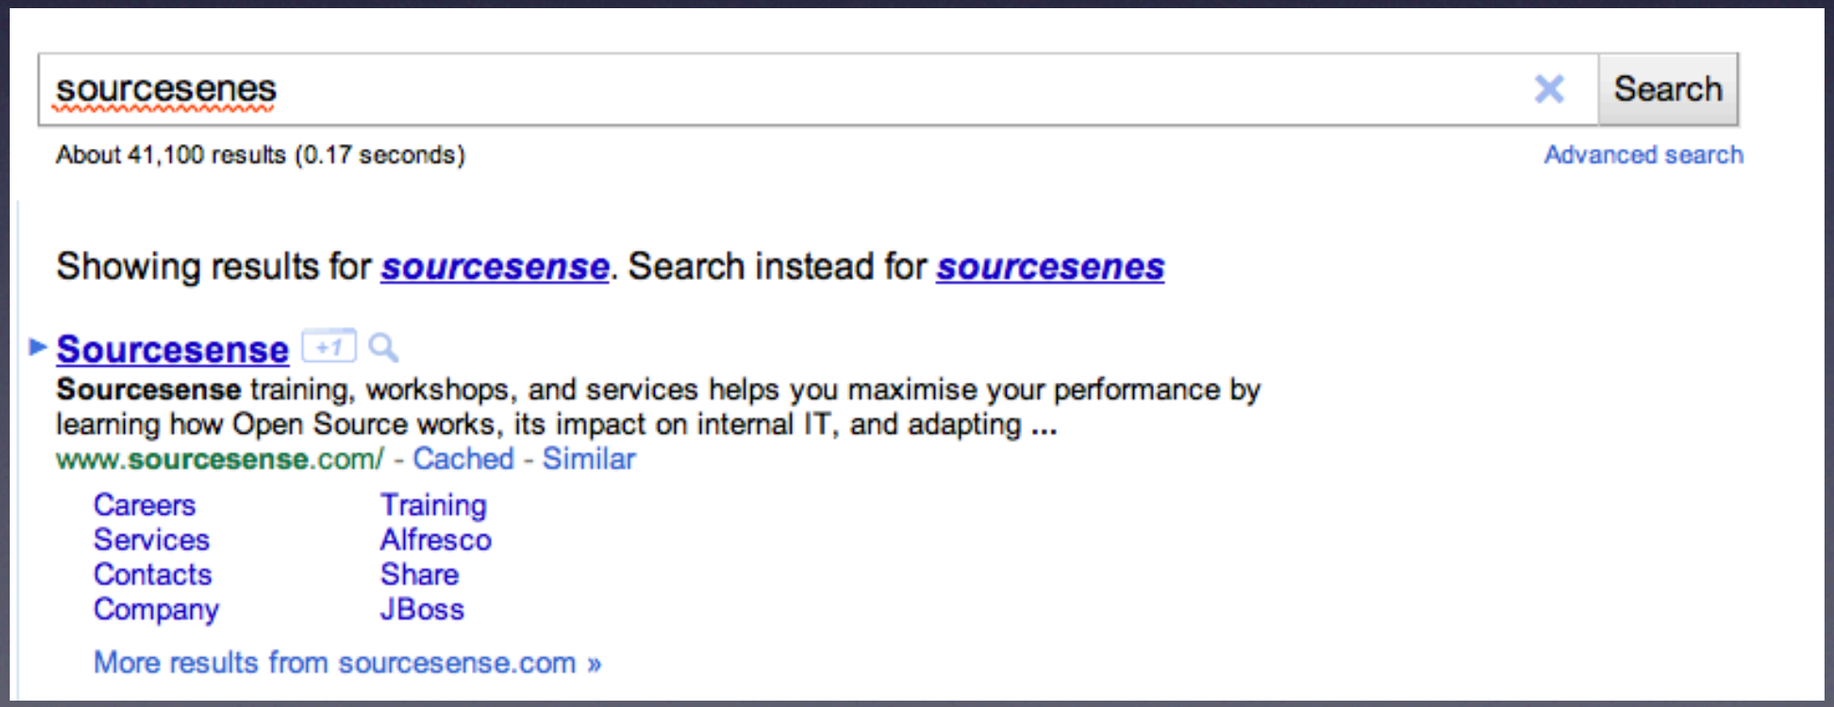

Solr - spellchecking

- build a spellcheck index dynamically
- return suggested results
- http://localhost:8983/solr/spell?q=hell ultrashar&spellcheck=true&spellcheck.collate=true &spellcheck.build=true
- useful to create custom query converters
  - <queryConverter name="queryConverter" class="org.apache.solr.spelling.SpellingQueryConverter"/>

Solr - similarity

- get documents “similar” to a given document or a set of documents
- Vector Space Model
- http://localhost:8983/solr/select?q=apache&mlt=true&mlt.fl=manu,cat&mlt.mindf=1&mlt.mintf=1&fl=id,score

Solr - geospatial search

- index location data
- query by spatial concepts and sort by distance
- find all documents with store position at no more than 5km than a specified point
- http://localhost:8983/solr/select?&indent=true&fl=name,store&q=*:*&fq={!geofilt%20sfield=store}&pt=45.15,-93.85&d=5

Solr - field collapsing

- group resulting documents on per field basis
  - http://localhost:8983/solr/select?&indent=true&fl=id,name&q=solr+memory&group=true&group.field=manu_exact
- useful for displaying results in a smart way
- see SOLR-236

Solr - join

- new feature (SOLR-2272)
- many users ask for it
- quite of a paradigm change
- http://localhost:8983/solr/select?q={!join+from=manu_id_s%20to=id}ipod&fl=id,manu_id&debugQuery=true

Solr Statistics: (example)

192.168.1.181

Category	[CORE] [CACHE] [QUERY] [UPDATE] [HIGHLIGHTING] [OTHER]
	Current Time: Fri May 27 09:06:53 CEST 2011
	Server Start Time: Fri May 27 09:00:12 CEST 2011

CORE

name:	`Searcher@5d352367 main`
class:	`org.apache.solr.search.SolrIndexSearcher`
version:	`1.0`
description:	`index searcher`
stats:	searcherName : Searcher@5d352367 main caching : true numDocs : 28 maxDoc : 28 reader : DirectoryReader(segments_2 _0(4.0):Cv28) readerDir : org.apache.lucene.store.NIOFSDirectory@/Users/tommasoteofili/Documents/workspaces/lucene_workspace/lucene_dev/solr/example/solr/data/index lockFactory=org.apache.lucene.store.NativeFSLockFactory@55f35e30 indexVersion : 1306332518058 openedAt : Fri May 27 09:00:12 CEST 2011 registeredAt : Fri May 27 09:00:12 CEST 2011 warmupTime : 0
name:	`core`
class:	
version:	`1.0`
description:	`SolrCore`
stats:	coreName : startTime : Fri May 27 09:00:12 CEST 2011 refCount : 2 aliases : []
name:	`searcher`

Solr statistics

evictions : 0
size : 0
warmupTime : 0
cumulative_lookups : 0
cumulative_hits : 0
cumulative_hitratio : 0.00
cumulative_inserts : 0
cumulative_evictions : 0

name:	`filterCache`
class:	`org.apache.solr.search.FastLRUCache`
version:	`1.0`
description:	`Concurrent LRU Cache(maxSize=4096, initialSize=1024, minSize=3686, acceptableSize=3891, regenerator=org.apache.solr.search.SolrIndexSearcher$2@358b6d8e)`
stats:	lookups : 129 hits : 118 hitratio : 0.91 inserts : 11 evictions : 0 size : 523 warmupTime : 11194 cumulative_lookups : 1889688 cumulative_hits : 1520086 cumulative_hitratio : 0.80 cumulative_inserts : 369603 cumulative_evictions : 0

Solr statistics

name:	/update
class:	org.apache.solr.handler.XmlUpdateRequestHandler
version:	$Revision: 1079955 $
description:	Add documents with XML
stats:	handlerStart : 1306248583062 requests : 253532 errors : 15 timeouts : 0 totalTime : 1649137 avgTimePerRequest : 6.50465 avgRequestsPerSecond : 1.0933348
name:	org.apache.solr.handler.FieldAnalysisRequestHandler
class:	org.apache.solr.handler.FieldAnalysisRequestHandler
version:	$Revision: 1065312 $
description:	Provide a breakdown of the analysis process of field/query text

Solr statistics

(WHAT IS THIS PAGE?)

Master	http://10.98.12.94:8681/solr/bd/replication
Poll Interval	00:05:00
Local Index	Index Version: 1294675577499, Generation: 208371
	Location: /mnt/LIVESOLR/slave/indexes/bd/TRLIVESOLR12/index.20110421120500
	Size: 21.04 GB
	Times Replicated Since Startup: 237
	Previous Replication Done At: Fri May 27 02:12:40 CEST 2011
	Config Files Replicated At: null
	Config Files Replicated: null
	Times Config Files Replicated Since Startup: null
	Next Replication Cycle At: Fri May 27 09:15:00 CEST 2011
Controls	Disable Poll
	Replicate Now
Cores:	[SHOP][DOTCOM][BD]
	Current Time: Fri May 27 09:11:23 CEST 2011
	Server Start At: Thu May 19 10:08:17 CEST 2011

Solr replication

Solr Architectures

- Simple
- Multicore
- Replication
- Sharded

Solr - MultiCore

- Define multiple Solr cores inside one only Solr instance
- Each cores maintain its own index
- Unified administration interface
- Runtime commands to create, reload, load, unload, delete, swap cores
- Cores can be thought as ‘collections’
- Allow no downtime while deploying new features/ bugfixes

Solr - Replication

- It's useful in case of high traffic to replicate a Solr instance and split (with eventually a VIP in front) the search load
- Master has the "original" index
- Slave polls master asking the latest version of index
- If slave has a different version of the index asks the master for the delta (rsync like)
- In the meanwhile indexes remain available
- No impact of indexing on search (almost)

Solr - Shards

- When an index is too large, in terms of space or memory required, it can be useful to define two or more **shards**
- A shard is a Solr instance and can be searched or indexed independently
- At the same time it's possible to query all the shards having the result be merged from the sub-results of each shard
- http://localhost:8983/solr/select?shards=localhost:8983/solr,localhost:7574/solr&q=category:information
- Note that the document distribution among indexes is up to the user (or who feeds the indexes)

Solr - Architectures

- When to use each?
- KISS principle
- High query load : replication
- Huge index : shard

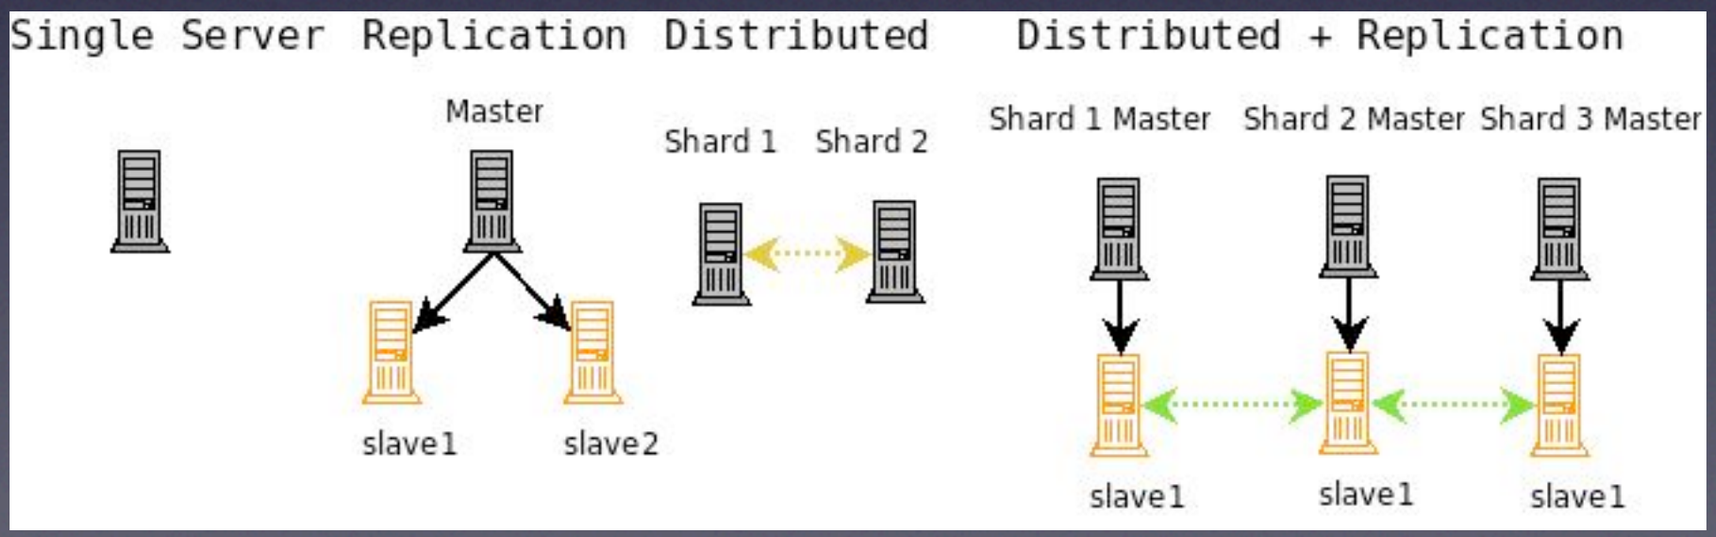

Solr - Architectures

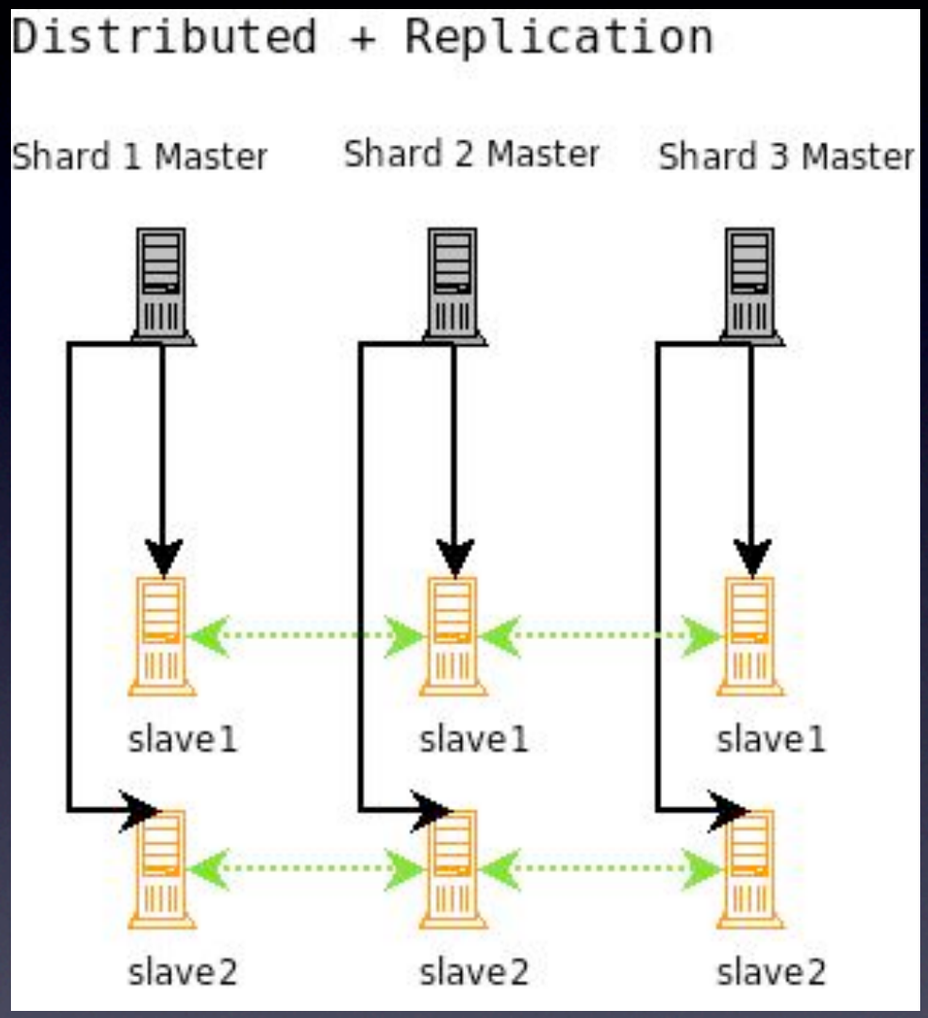

- High queries w/ large indexes : shard + replication

Solr - Architectures

- Tips & Tricks:
  - Don’t share indexes between Master and Slaves on distributed file systems (locking)
  - Anyway get rid of distributed file systems (slow)
  - Lucene/Solr is I/O intensive thus behaves better with quick disks
  - Always use MultiCore - hot deploy of changes/ bugfixes
  - Replication is network intensive
  - Check replication poll time and indexing rates

Tips&Tricks

- Solr based SE development process
- Plugins
- Performance tuning
- Deploy

Process - t_0 analysis

- Analyze content
- Analyze queries
- Analyze collections
- Pre-existing query/index load (if any)
- Expected query/index load
- Desired throughput/avg response time
- First architecture

Process - n-th iteration

- index 10-15% content
- search stress test (analyze peaks) - use SolrMeter
- quality tests from stakeholders (accuracy, recall)
- eventually add/reconfigure features
- check http://wiki.apache.org/solr/FieldOptionsByUseCase and make sure fields used for faceting/sorting/highlighting/etc. have proper options
- need to change field types/analysis/options - rebuild the index

Solr - Plugins

- QParserPlugin
- RequestHandler (Search/UpdateHandler)
- UpdateRequestProcessor
- ResponseWriter
- Cache

Performance tuning

- A huge tuning is done in schema.xml
- Configure Solr caches
- Set auto commit where possible
- Play with mergeFactor

Performance tuning

- The number of indexed fields greatly increases memory usage during indexing, segment merge time, optimization times, index size
- Stored fields impact on index size, search time, ...
- set omitNorms=”true” where it makes sense (disabling length normalization and index time boosting)
- set omitTermFreqAndPositions=”true” if no queries on this field using positions or should not influence score

Performance tuning

- FilterCache - unordered document ids for caching filter queries
- QueryResultCache - ordered document ids for caching queries results (caching only the returned docs)
- DocumentCache - stores stored fields (at least <max_results> * <max_concurrent_queries>
- Setup autowarming - keep caches warm after commits

Performance tuning

- Choose correct cache implementation FastLRUCache vs LRUCache
- FastLRUCache has faster gets and slower puts in single threaded operation and thus is generally faster than LRUCache when the hit ratio of the cache is high (> 75%)

Performance tuning

- Explicit warm sorted fields
- Often check cache statistics
- JVM options - don’t let the OS without memory!
- mergeFactor - impacts on the number of index segments created on the disk
  - low mF : smaller number of index files, which speeds up searching but more segment merges slow down indexing
  - high mF : generally improves indexing speed but gets less frequent merges, resulting in a collection with more index files which may slow searching

Performance tuning

- set autocommit where possible, this will avoid close and reopen of IndexReaders everytime a document is indexed - can choose max number of documents and/or time to wait before automatically do the commit
- finally...need to get your hand dirty!

Deploy

- SolrPackager by Simone Tripodi!
- It's a Maven archetype
- Create standalone/multicore project
- Each project will generate a master and a slave instance
- Define environment dependent properties without having to manage N config files
- 'mvn -Pdev package' // will create a Tomcat package for the development environment

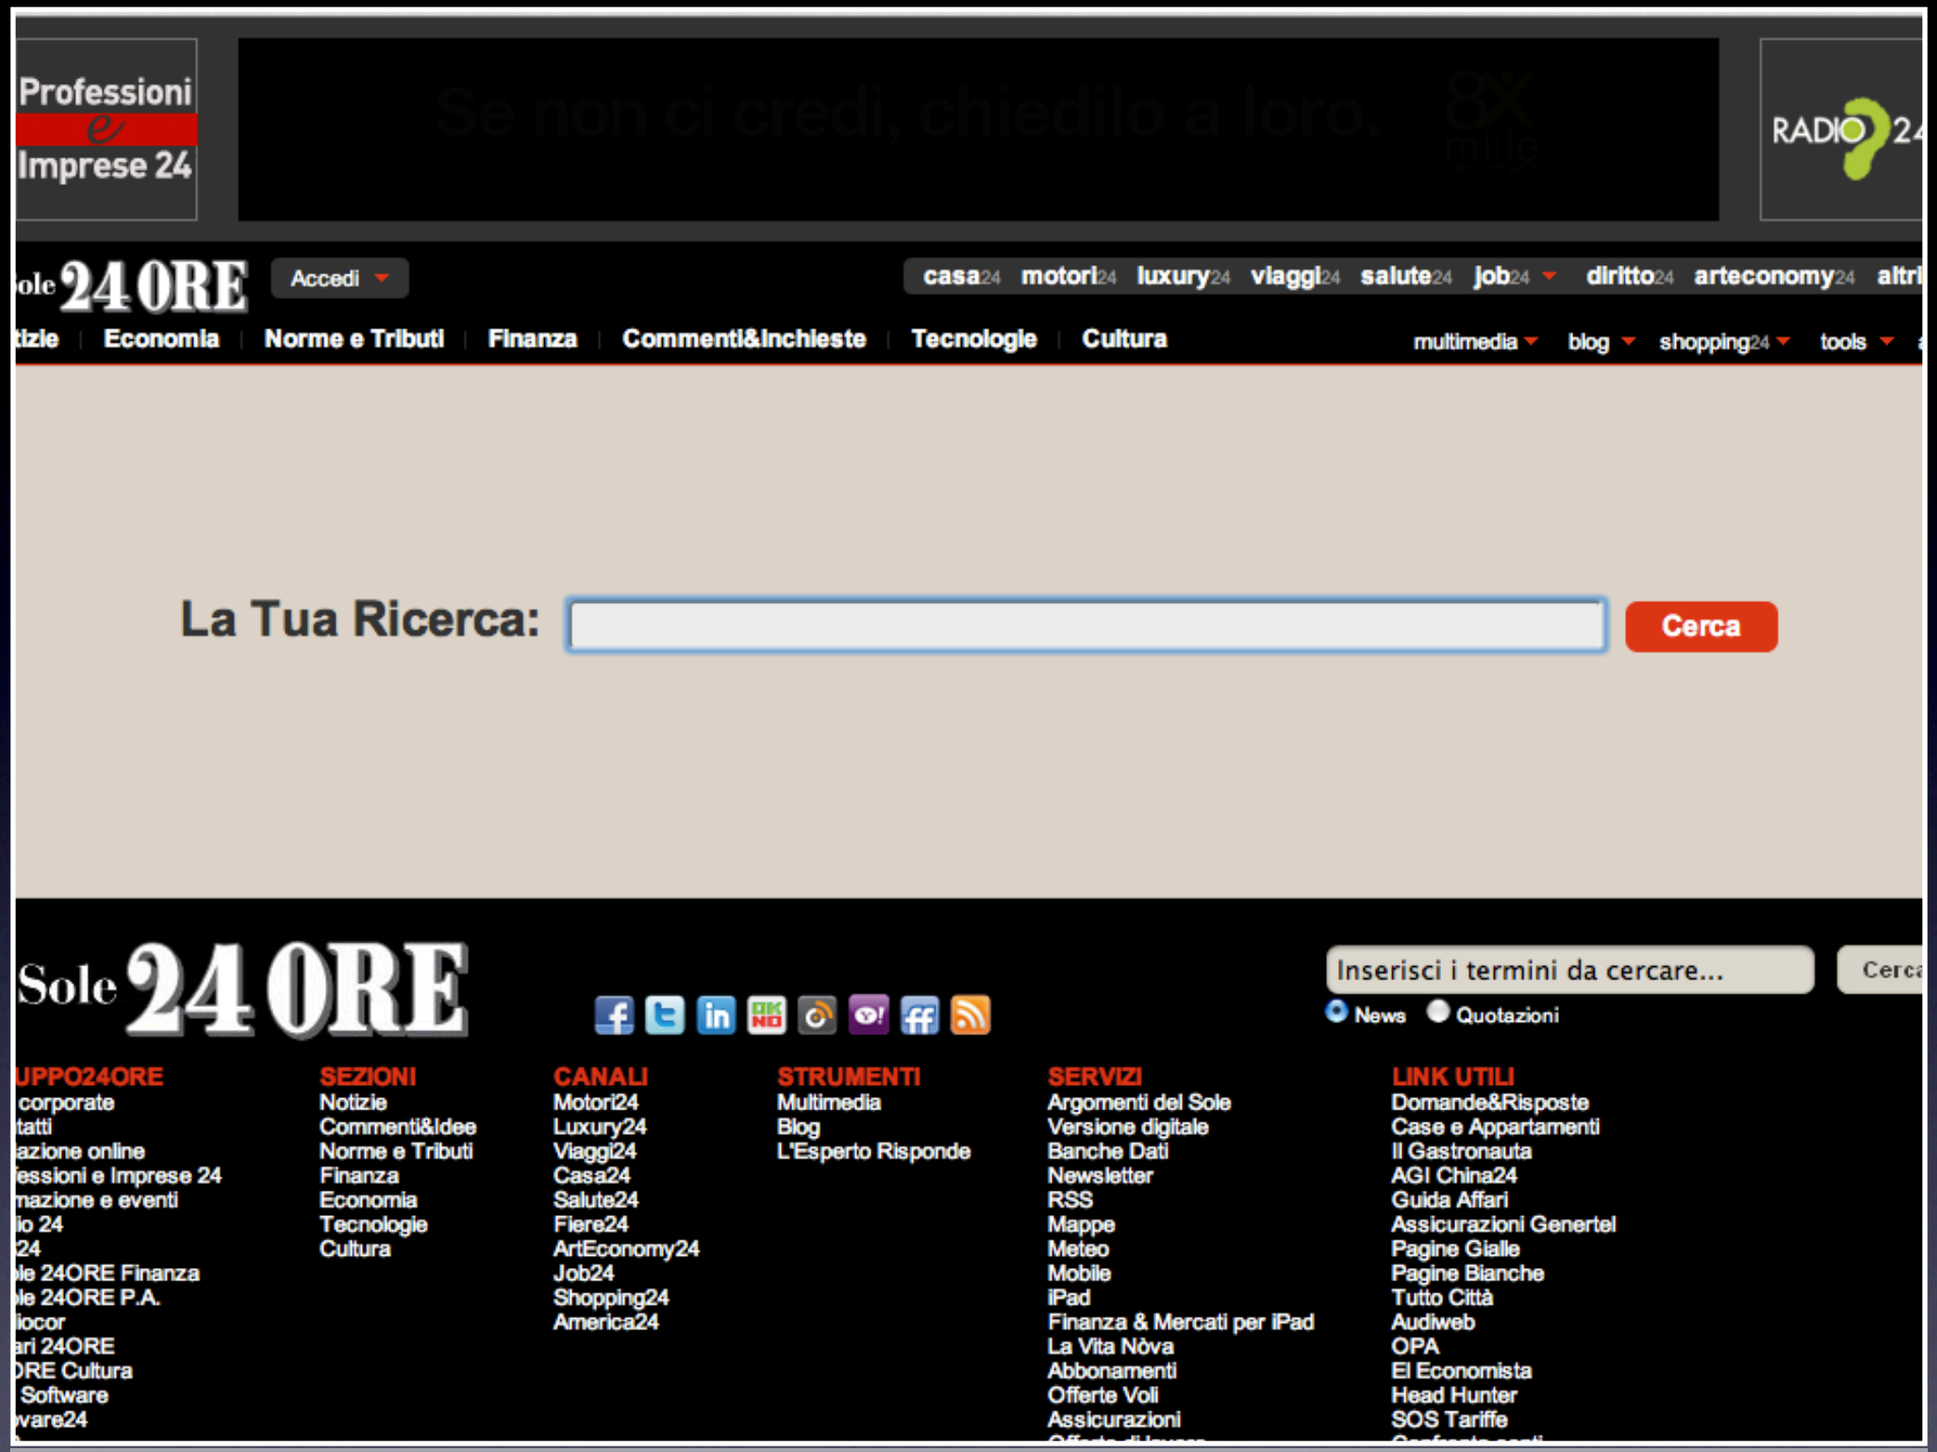
Professioni e Imprese 24
Se non ci credi, chiedilo a loro.
8x mille
RADIO 24
Sole 24 ORE
Accedi
casa24 motori24 luxury24 viaggi24 salute24 job24 diritto24 arteconomy24 altri
Economia | Norme e Tributi | Finanza | Commenti&Inchieste | Tecnologie | Cultura
multimedia blog shopping24 tools
La Tua Ricerca:
Cerca
Sole 24 ORE
Inserisci i termini da cercare...
Cerca
News Quotazioni
SEZIONI
Notizie
Commenti&Idee
Norme e Tributi
Finanza
Economia
Tecnologie
Cultura
CANALI
Motori24
Luxury24
Viaggi24
Casa24
Salute24
Fiere24
ArtEconomy24
Job24
Shopping24
America24
STRUMENTI
Multimedia
Blog
L'Esperto Risponde
SERVIZI
Argomenti del Sole
Versione digitale
Banche Dati
Newsletter
RSS
Mappe
Meteo
Mobile
iPad
Finanza & Mercati per iPad
La Vita Nòva
Abbonamenti
Offerte Voli
Assicurazioni
LINK UTILI
Domande&Risposte
Case e Appartamenti
Il Gastronauta
AGI China24
Guida Affari
Assicurazioni Genertel
Pagine Gialle
Pagine Bianche
Tutto Città
Audiweb
OPA
El Economista
Head Hunter
SOS Tariffe

Case study

La Tua Ricerca : francesco totti

Cerca

Risultati **1-10** su **1465** per la ricerca **" francesco totti "** ordina: per rilevanza filt

Francesco Totti

Tutte le notizie e le informazioni su Francesco Totti nella scheda degli **Argomenti del Sole**.

http://argomenti.ilsole24ore.com/francesco-totti.html

Adesso **Totti** sogna di superare Nordhal

di Gigi Garanzini- 03/05/2011

Adesso **Totti** sogna di superare Nordhal ...il sorpasso di **Totti** su Roberto Baggio nella classifica cannonieri di tutti i tempi.

http://www.ilsole24ore.com/art/notizie/2011-05-03/adesso-totti-sogna-superare-084816.shtml

Totti supera Baggio tra i cannonieri di tutti i tempi, Milan a un passo dallo scudetto

di Gigi Garanzini- 02/05/2011

Totti supera Baggio tra i cannonieri di tutti i tempi, Milan a un passo dallo scudetto Un primo maggio di grandi numeri, se è vero che con la doppietta di Bari capitan **Totti** ha sorpassato Baggio tra i cannonieri di tutti i tempi,...

http://www.ilsole24ore.com/art/notizie/2011-05-02/commento-maggio-082732.shtml

Risultati Multimediali per "**francesco totti**"

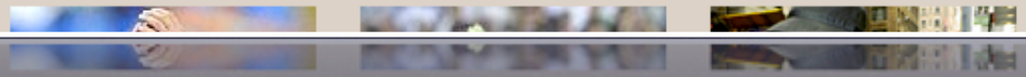

AFFINA LA RICERCA:

- Roma (225)
- capitano (67)
- squadra (64)
- giocatore (52)
- Inter (44)
- gol (41)
- derby (40)
- Milan (34)
- nazionale (34)
- Luciano Spalletti (32)

RISULTATI DI:

- Oggi

Risultati da **Shopping24**

€ 13
Iva In

Francesco totti

Vita, parole e impres
dell'ultimo gladiatore
di Cecchini Massimo

€ 18
Iva In

Tutto totti

Tutte le barzellette su
nuove barzellette su
mo je faccio er cucch
Cofanetto con tre vol
indivisibili
di [Totti Francesco, F

Case study

Case Study

- Architecture analysis
- Plugin development
- Testing and support

Challenges

- Architecture
- Schema design

Challenge

- Architecture
  - 4B docs of ~4k each
  - ~3 req/sec overall
- 3 collections:
  - |archive| = 3B
  - |2010-2011| = 1M
  - |intranet| = 0.9B

Challenge

- Content analysis
- get the example Solr schema.xml
- optimize the schema in order to enable both stemmed and unstemmed versions of fields: author, title, text, cat
- add omitNorms=”true” where possible
- add a field ‘html_content’ which will contain an HTML text but will be searched as clean text
- all string fields should be lowercased

Extras

- Clustering (Solr-Carrot2)
- Named entity extraction (Solr-UIMA)
- SolrCloud (Solr-Zookeeper)
- ManifoldCF
- Stanbol EntityHub
- Solandra (Solr-Cassandra)

THANKS!